

Git Tutorial Tour

Jon Loeliger



Git Tutorial Tour Outline

- Repository Management
 - Git Repository Concepts
 - Inter-Repository Data Flow
 - Branches
 - Merges
- Local Use Cases
 - Pull Updates
 - Generating Patches
 - Examples
- Supplemental Slides

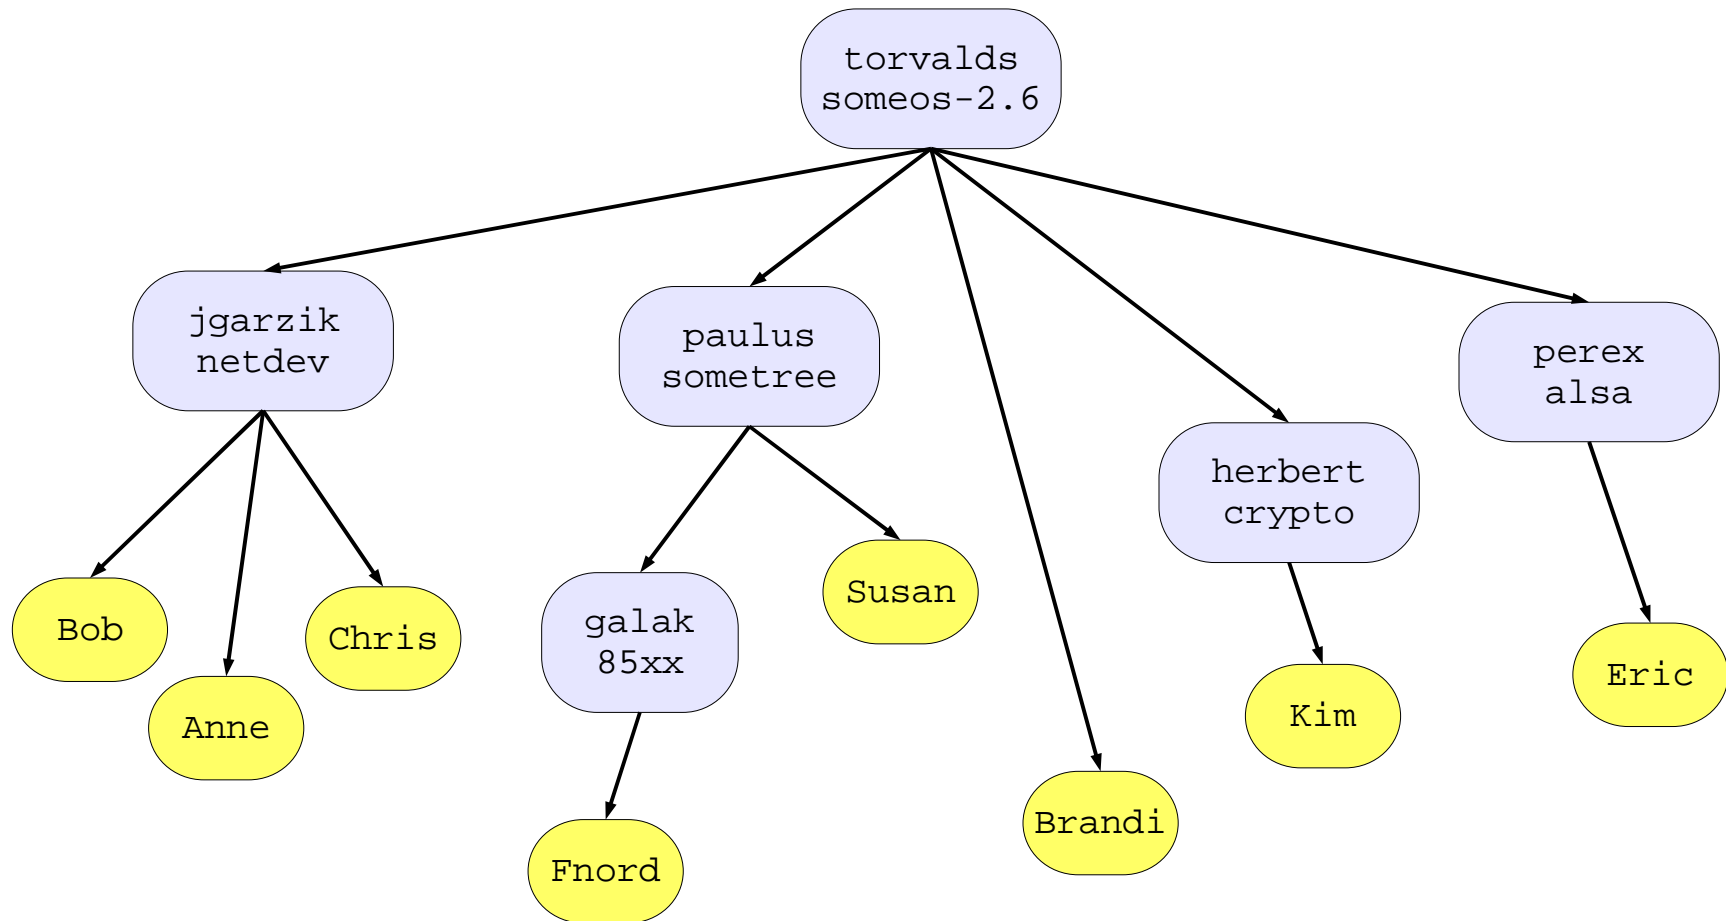
Git Concepts: Repositories

- Git repository is self-contained, local and complete
 - Doesn't require external content
 - No centralized repository
 - Includes complete history of every file and change log
- Git repository is not distributed
 - No: part here, part there, part over there...
- No git repository is inherently authoritative
 - Only authoritative by convention and agreement

Git Concepts: Distributed Development

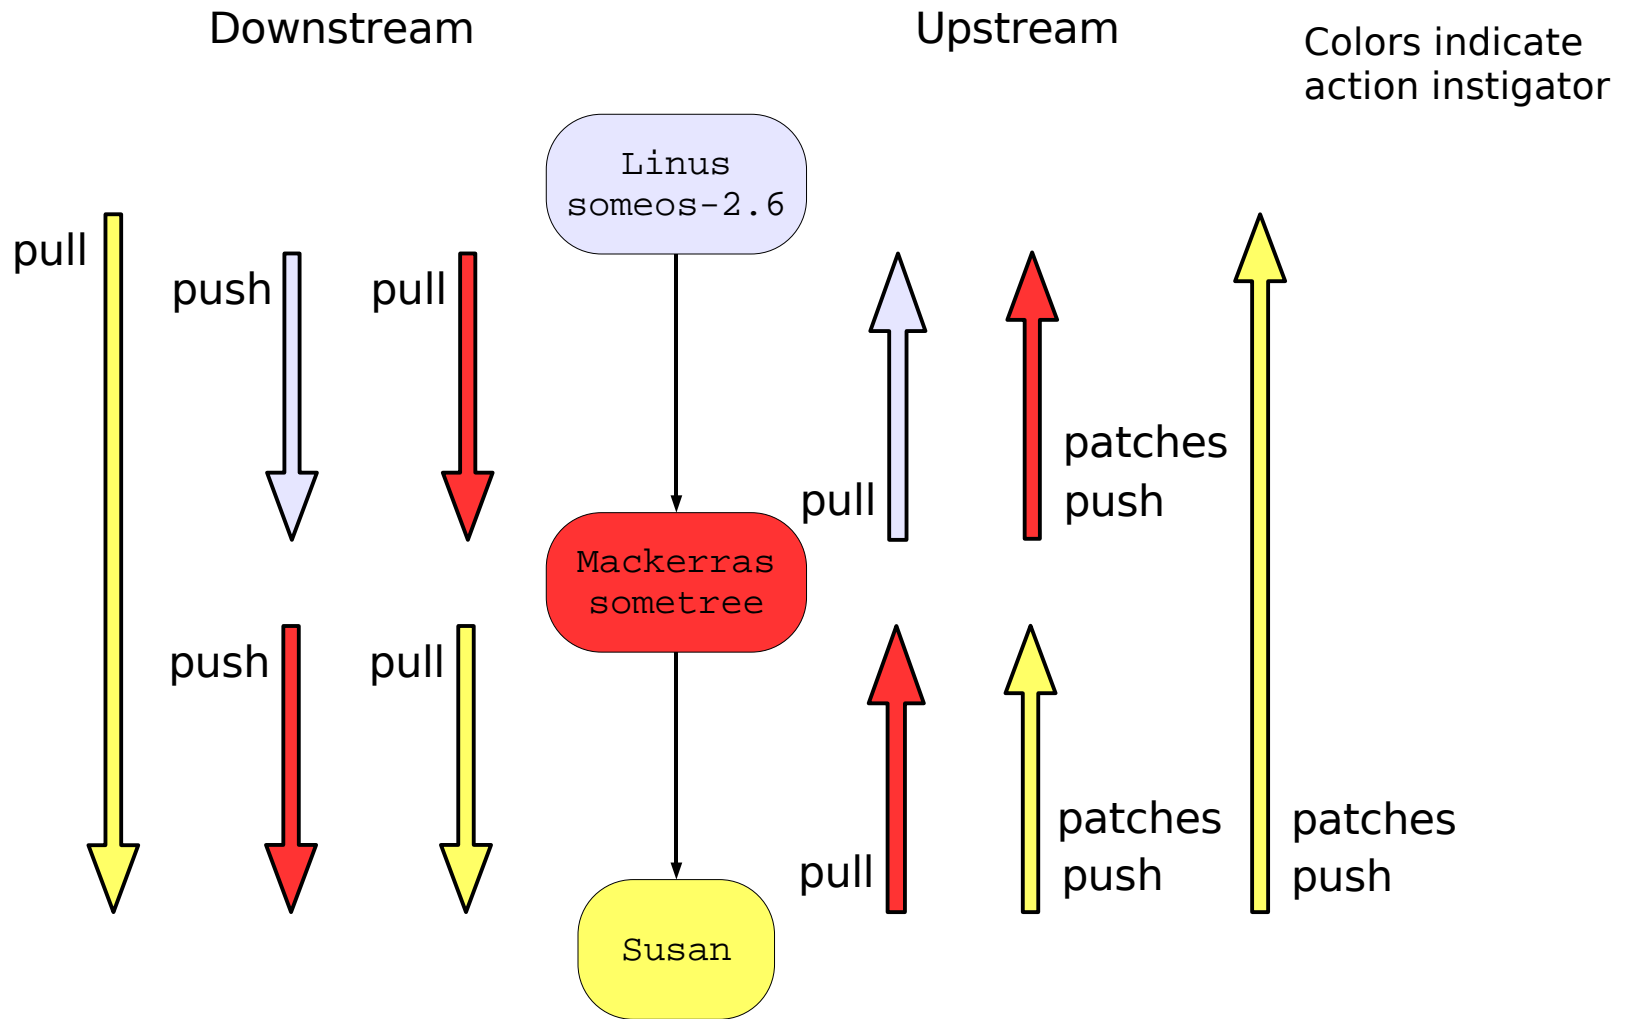
- Development can be distributed
 - Multiple repositories can create different histories and changes
 - ◆ Even if they originated from a common ancestor
 - Multiple repositories with different development can be combined
- Development can be shared
 - Multiple developers can use and update a common repository
 - ◆ No matter if the repository is local or remote
- Development can be private
 - Revision control your Address List at home

Repository Network



A → B B is cloned from A

Repository Data Flow



Git Concepts: Branches in git

- Cheap, fast and easy
- Topic/Development Branches
 - Stable, development, bug-fix, testing branches
 - Small development lines, per-feature, per-developer, etc
 - Collect, reorder or organize changes
 - Cherry-pick particular patches
- Tracking Branches
 - Follow upstream changes in local repository
 - Don't commit to tracking branches
 - Identified as RHS of `PULL : refsspecs`
- Branch names refer to the current branch HEAD revision
- Branches don't have a “beginning” per se
 - `$ git merge-base <original-branch> <branch-name>`


Branches: Cloning

- Clone from upstream URL

```
$ git clone  
git://www.jdl.com/software/dtc.git
```


- Copies complete upstream repository into a local repository
- Creates `origin` tracking branch
 - Tracks upstream `master` branch
- Introduces `master` development branch
 - Initially the same place as `origin`
 - For *your* development

Original Repository



git clone

Cloned Repository



Branches: Commit Onto Branches

- Master branch in your repo is for *your* development!

```
$ git checkout master  
# Edit a file.  
$ git commit -a -m "Add copyright.  Fix 80-column line."
```

- Introduce your own topic development branches too

```
$ git checkout -b jdl  
# Edit a file.  
$ git commit -a -m "Remove dead code."
```

- Gotcha: Commits require an author!

- Environment variables: `GIT_AUTHOR_NAME`, `GIT_AUTHOR_EMAIL`
- Config file: `$ git repo-config user.name 'Eric S. Raymond'`
- Gecos value: `/etc/passwd`

Branches: Visualizing Branches

```
$ git show-branch
```

```
! [jdl] Remove dead code.
```

```
* [master] Add copyright. Fix 80-column line.
```


```
! [origin] dtc: add setting of physical boot cpu
```

```
---
```


```
* [master] Add copyright. Fix 80-column line.
```

```
+ [jdl] Remove dead code.
```

```
+*+ [origin] dtc: add setting of physical boot cpu
```



Branch Timeline



Git Concepts: Merges

- Merge: Combine directory and file contents from separate sources to yield one combined result.
 - Sources for merges are local branches
 - Merges always occur in the current, checked-out branch
 - A complete merge ends with a new commit
- Merge resolution is inherently ambiguous
 - Git uses several merge heuristics:
 - ♦ Several merge strategies: resolve, recursive, octopus, ours
 - ♦ Techniques: fast-forward, three-way
- Some merge conflicts may need to be resolved by the developer

Merge jdl branch To master


- Merge jdl branch into master

```
$ git checkout master  
$ git pull . jdl
```

- Resulting Branch History

```
$ git show-branch
```

```
! [jdl] Remove dead code.  
* [master] Merge branch 'jdl'  
! [origin] dtc: add setting of physical boot cpu  
---  
- [master] Merge branch 'jdl'  
+* [jdl] Remove dead code.  
* [master^] Add copyright. Fix 80-column line.  
+*+ [origin] dtc: add setting of physical boot cpu
```



Outline: Use Cases

- Repository Management
- Local Use Cases
 - Pulling Updates
 - Making Patches
 - Finding a Commit
 - Cherry Picking
 - Reverting a Commit
 - Resolving Merges
 - Rebasing Local Changes
- Supplemental Slides

Pulling Updates

- The default origin “remotes” file

```
$ cat .git/remotes/origin
URL: git://www.jdl.com/software/dtc.git
Pull: refs/heads/master:refs/heads/origin
```


- Get updates from remote repository

```
$ git checkout master
$ git pull origin
```

- Pull is a fetch then merge


- Places fetched updates in `ref/heads/origin`
- Merges `origin` into the current, checked-out branch, `master`

Original Repository



git pull


Cloned Repository



Pull: Fetch and Merge Branch Timeline


Before

Public Repository




Your Repository

time




After

Public Repository



1- fetch

2- fast-forward



3 - merge

Your Repository

time



Pull: What Else?

- Pull is a `fetch` followed by a `merge`
 - Merges should start with some common base, though
- Can fetch/pull any branch from any repository
 - Supply URL on command line directly
 - Maybe add new “remotes” file with URL and `Pull:` line
 - Even this repository: `git pull . <branch>`
- Could just `git fetch` an upstream branch too
 - Able to then `cherry-pick` commits rather than full merge
 - Able to use `git diff`, `git log`, etc.

Pull: When?

- When you want to!
- When it is convenient, stable, Wednesday...
- When the upstream tree is stable
- When your working directory is stable
 - ...and you have the right branch checked out
 - Your working directory should probably be clean
 - ♦ `git diff` or `git status` should be empty
 - Technically not necessary, but...

Sending Changes Upstream

- Generate and send patches via email
 - Most developers send patches to a maintainer or list
 - Highly visible public review of patches on mail list
- Maintainer pulls updates from a downstream developer
 - Maintainer can directly pull from your published repository
 - Initiated by upstream maintainer
- Developer pushes updates to an upstream maintainer
 - Some developers have write permissions on an upstream repository
 - Initiated by downstream developer

Patches: Generate a Patch

- Use `git format-patch` command

```
$ git format-patch --signoff origin..jdl
```

```
From: Jon Loeliger <jdl@jdl.com>
```

```
Date: Sat, 24 Jun 2006 15:42:51 -0500
```

```
Subject: [PATCH] Remove dead code.
```

```
Signed-off-by: Jon Loeliger <jdl@jdl.com>
```

```
---
```

```
data.c | 18 -----
```

```
1 files changed, 0 insertions(+), 18 deletions(-)
```

```
diff --git a/data.c b/data.c
```

```
index 911b271..d3b55f9 100644
```

```
--- a/data.c
```

```
+++ b/data.c
```

```
@@ -20,24 +20,6 @@
```

Patches: A Side Note

- Read the README
- Follow the coding standards
- Compile and test your code
- Fix the whitespace issues
- Know who the upstream maintainer is
- Understand any “sign-off” policy

Patches: Formatting and Sending Mail

- Send plain ASCII, not HTML
 - Send patches inline
 - Don't use attachments
- Send only your changes
- One line “commit summary” first
- Be careful of cut-n-paste solutions
 - Use file browse and insert if possible
- Maybe move pre-formatted patches to your MUA draft folder
- Send git-format-patch output directly

```
$ git send-email --to  
maintainer@example.com my.patch
```

Finding a Commit

- SHA-1 hash names are constant
- Symbolic commit names like `master~4` change over time!
 - They are relative to the current HEAD
- Ways to determine the name of a commit you want
 - `git show-branch => paul~38^2~10^2`
 - Paw around in `gitk => 9ad494f62444ee37209e85173377c67612e66ef1`
 - Use `git log -- <some/file/path>`
 - Use `gitweb`
- Note:

```
$ git rev-parse paul~38^2~10^2
9ad494f62444ee37209e85173377c67612e66ef1
```

Cherry Picking Example

- You want one or more particular commits from some other branch applied to your current branch
 - Bug fix from some other branch
 - Transfer partial functionality from development branch
- Get that commit into your repository
 - Already present on a different branch
 - Perhaps using `git fetch` from another repository
- Cherry-pick it into appropriate branch


```
$ git checkout master
$ git cherry-pick 9ad494
```

Cherry Picking: Before and After


- Branch Timeline Picture

- 4' is a different commit with the same content as 4

Before



After



Revert Example

- Situation:
 - Some commit buried in the past needs to be undone
 - ... and it has already been merged to master too!

```
$ git show-branch --more=10 master mpc86xx
```


```
* [master] Merge branch 'mpc86xx'
! [mpc86xx] Remove obsolete #include <someos/config.h>.
--
- [master] Merge branch 'mpc86xx'
*+ [mpc86xx] Remove obsolete #include <someos/config.h>.
*+ [mpc86xx^] Remove redundant STD_MMU selection.
*+ [mpc86xx~2] Move I8259 selection under MPC8641HPCN board.
*+ [mpc86xx~3] Remove redundant PPC_86XX check.
*+ [mpc86xx~4] Reworked the IRQ mapping. Now reading IRQ and ... <===== ARGH!
*+ [mpc86xx~5] Make 86xx secondary CPU start be more generic.
*+ [mpc86xx~6] Updated 86xx defcnofig after merge with Paul
```

- This is a job for `git revert`


Revert: Before and After

- Commit 4 is reverted, creating commit 0 in yellow
- Commit 0 is merged to master as commit 0'

Before



After



Revert: Commands

- Revert applies a new “inverse” patch for a given commit
- Revert works without rewriting the commit history
- Can revert multiple commits at the same time

- Revert commit on `mpc86xx` topic branch:

```
$ git checkout mpc86xx  
$ git revert mpc86xx~4
```

- Reflect it onto `master` branch as well:

```
$ git checkout master  
$ git pull . mpc86xx
```

Revert: Branch History After

```
$ git show-branch --more=10 master mpc86xx

* [master] Merge branch 'mpc86xx'
! [mpc86xx] Revert "Reworked the IRQ mapping. Now reading IRQ and ..."
--
- [master] Merge branch 'mpc86xx'
*+ [mpc86xx] Revert "Reworked the IRQ mapping. Now reading IRQ and ..."
*+ [mpc86xx^] Remove obsolete #include <someos/config.h>.
*+ [mpc86xx~2] Remove redundant STD_MMU selection.
*+ [mpc86xx~3] Move I8259 selection under MPC8641HPCN board.
*+ [mpc86xx~4] Remove redundant PPC_86XX check.
*+ [mpc86xx~5] Reworked the IRQ mapping. Now reading IRQ and ...
*+ [mpc86xx~6] Make 86xx secondary CPU start be more generic.
*+ [mpc86xx~7] Updated 86xx defcnofig after merge with Paul
```

Revert versus Reset versus Checkout

- Reset versus Revert: Trying to undo the last commit?
 - Ask: Does someone else have this version of the repository history?
 - Yes: Use `git revert` and do not change history
 - No : Could use `git reset` or `git commit --amend` perhaps
- Reset versus Checkout
 - Reset doesn't change your current branch
 - Checkout establishes your current branch
 - These can be similar unless naming a different branch:

```
git reset --hard topic    # still on original branch,  
                          #   but it now looks like topic!  
git checkout -f topic     # topic branch checked out now
```

Merge Resolution Hell ... er, Examples

- Pulling updates from upstream can go wrong for a number of reasons!
 - Conflicting file contents
 - New or modified files in your working tree
 - Different file modes
 - Possible for your upstream changes to come back differently
 - ♦ Altered upstream by maintainer?
 - ♦ Were they pulled or patched upstream?
 - ♦ Same file *content*, but different *commits*?
 - Merge is a heuristic!

Merge Resolution Strategies

- Before doing the pull or merge

- Check out the correct “merge into” branch, likely `master`
- Ensure a clean working directory first

```
git ls-files --others
git status
```

- After “failed” pull/merge request:

```
git ls-files -u    # Show unmerged files remaining
git diff          # Clean if it matches one of the variants!
```

- Make progress resolving conflicts:

```
git update-index  # Tell git when a conflict has been resolved
```

- Be done :

- Fully resolved: `git commit`
- Abandon merge: `git reset --hard ORIG_HEAD`

Failed Merge: Untracked Working Tree File

- Git says:

```
git-read-tree: fatal: Untracked working tree file  
'Documentation/ABI/README' would be overwritten by merge.
```

- Huh? I don't have that file...

- Maybe you *do* have that file and should check!
- But maybe it is leftover from a previous merge effort?
- Are you on the right branch?

```
git branch
```

```
git show-branch
```


Failed Merge: Untracked File Resolution

- Your key:

```
git ls-files --others    # Look for untracked files
git status              # Check for unexpected files
git clean -d            # Remove cruft
```

- Reset the working directory:

- Maybe abandon this merge?
- Look for other files
- Clean out the old files
- Do the merge again!

```
$ git reset --hard ORIG_HEAD
$ git ls-files --others
$ git clean -d
$ git status    # Ah! Nice-n-tidy
```

Failed Merge: CONFLICT Content

- Git message:

CONFLICT (content): Merge conflict in drivers/net/phy/Makefile

- Semi-traditional "<<<< ==== >>>>" style file content differences.

```
$ git diff drivers/net/phy/Makefile
```

```
@@@ -8,4 -8,5 +8,9 @@@
```

```
obj-$(CONFIG_CICADA_PHY)      += cicada.o
```

```
obj-$(CONFIG_LXT_PHY)        += lxt.o
```

```
obj-$(CONFIG_QSEMI_PHY)      += qsemi.o
```

```
++<<<<<< HEAD/drivers/net/phy/Makefile
```

```
+obj-$(CONFIG_VITESSE_PHY)    += vitesse.o
```

```
++=====
```

```
+ obj-$(CONFIG_SMSC_PHY)      += smsc.o
```

```
+ obj-$(CONFIG_VITESSE_PHY)  += vitesse.o
```

```
++>>>>>>
```

```
501b7c77de3e90519e95fd99e923bf9a29cd120d/drivers/net/phy/Makefile
```

Failed Merge: Resolve Content Conflict

- Your key:

```
git diff
```

```
git ls-files -s arch/somearch/Kconfig    # Various stage versions
```

- Edit and fix

- Use your favorite editor, emacs

- Make progress

```
git update-index arch/somearch/Kconfig
```

Failed Merge: CONFLICT add/add

- Git says either:
 - Added file3 in both, but differently. ERROR: Merge conflict in file3
 - CONFLICT (add/add):
File arch/somearch/platforms/86xx/mpc86xx_hpcn.c
added non-identically in both branches.
Adding as arch/somearch/platforms/86xx/mpc86xx_hpcn.c~HEAD
and arch/somearch/platforms/86xx/mpc86xx_hpcn.c~501b7c instead.
- Two branches added the same file, but differently.
 - The version from my master HEAD is left with ~HEAD suffix.
 - The version from incoming origin is left with ~501b7c suffix.
- Resolution:
 - Pick the right one or merge them
 - Use `git add` to add it back with the right name (drop ~suffix)
 - Remove the other

Failed Merge: CONFLICT rename/add

- Git says:

```
CONFLICT (rename/add): Rename  
arch/{mips/configs/ddb5476_defconfig =>  
  somearch/configs/mpc8641_hpcn_defconfig} in 501b7c...  
arch/somearch/configs/mpc8641_hpcn_defconfig added in HEAD
```

- Maybe it lies!

- Rename detection isn't 100% precise. It is a heuristic!
 - ddb5476_defconfig just happens to be removed
 - mpc8641_hpcn_defconfig really added!

- Resolution:

- Add it back! Holy cow!
- Make sure index has correct, added file and either:

```
git update-index arch/somearch/configs/mpc8641_hpcn_defconfig  
git add arch/somearch/configs/mpc8641_hpcn_defconfig
```


Rebase Local Changes Example

- Situation:
 - You have done development work on your topic branch
 - You pull in upstream `origin` and merge it to `master`
 - You don't want to merge `master` into your topic branch
 - You don't want to merge your topic into the `master` branch
 - No one else has cloned your topic branch!
 - ♦ Why? Don't rewrite history
 - But you want to send your topic branch patches upstream!
- This is a job for `git rebase`


Rebase: Before and After

- Rebase commits 0 – 4 of `topic` onto `master` branch at A as 0' – 4'

Before



After



Rebase: Commands

- Rebase Commands

```
$ git checkout topic  
$ git rebase master
```

- Creates a series of patches from `topic` to be applied to `master`

- May have to resolve conflicts at each step due to merge operation

```
git rebase --continue  
git rebase --skip  
git rebase --abort
```


Git Resources

- Sources and Documentation:
 - Git sources, documentation and many repositories: kernel.org/git
 - The Git Wiki: git.or.cz/gitwiki
- Front-ends and Viewers:
 - Cogito: kernel.org/pub/software/scm/cogito
 - Stacked Git: www.procode.org/stgit
 - Patchy Git: www.spearce.org/category/projects/scm/pg
 - QT Gui viewer: sourceforge.net/project/qgit
- Mail List: [git @ vger.kernel.org](mailto:git@vger.kernel.org)
- IRC: [#git](http://freenode.net)

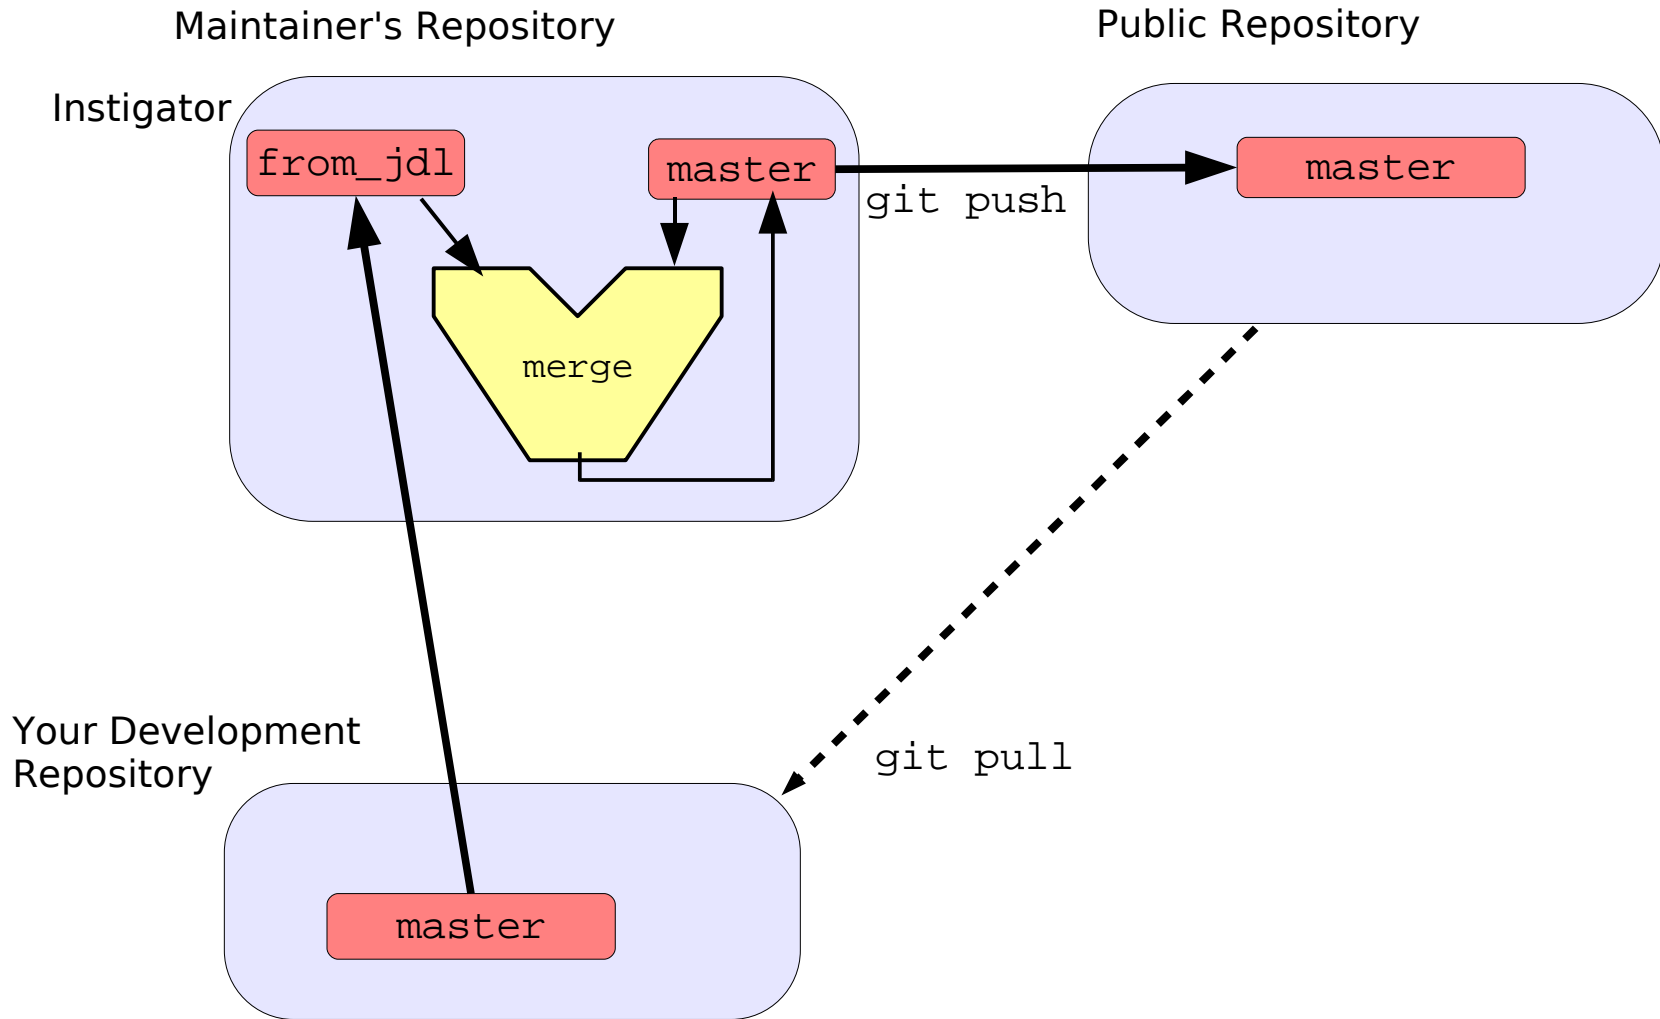
Outline: Supplemental

- Repository Management
- Local Use Cases
- Supplemental Slides
 - Sending Changes Upstream Using Pull and Push
 - Screen Capture of gitk

Send Changes Upstream Using git pull

- Upstream maintainer trusts you and your work
- You have the ability to publish a public repository
- You have a lot of changes, more than a few patches
- Advertise your repository and branch
- Wait for upstream maintainer to pull it
 - `$ git pull git://www.your-site.org/path/to/repo.git`
- Profit.

Using Pull: Data Flow Picture



Sending Changes Upstream Using git push

- Direct push into remote repository
 - Sends objects, packs, branches to remote repository
- Push to a repository from which you have fetched
 - Can only push to a branch that is a proper subset
 - Should result in a "fast-forward" on the remote end
 - Technically you can push elsewhere, but...
- Requires write access on the remote end
 - Likely via ssh

Using Push: Setup Remotes File

- Create a “remotes” file
 - State upstream repository URL
 - State branches to be pushed upstream
- Can push multiple branches, as needed
 - Just add a `Push: refsPEC` line for each branch
- Can push to different remote branch names

`Push: master:incoming`

```
$ cat .git/remotes/publish
URL: ssh://www.jdl.com/software/dtc.git
Push: master:master
Push: jdl:jdl
```

Using Push: git push

```
$ git push publish
```

Password:

```
updating 'refs/heads/master'
```

```
  from 38e8f8fd88dae07ef8ada9d6baa41b06a4d9ac9f
```

```
  to   a73b7d43d4f60e76d82018fb9a4d137b089a1325
```

```
Generating pack...
```

```
Done counting 11 objects.
```

```
Result has 8 objects.
```

```
Deltifying 8 objects.
```

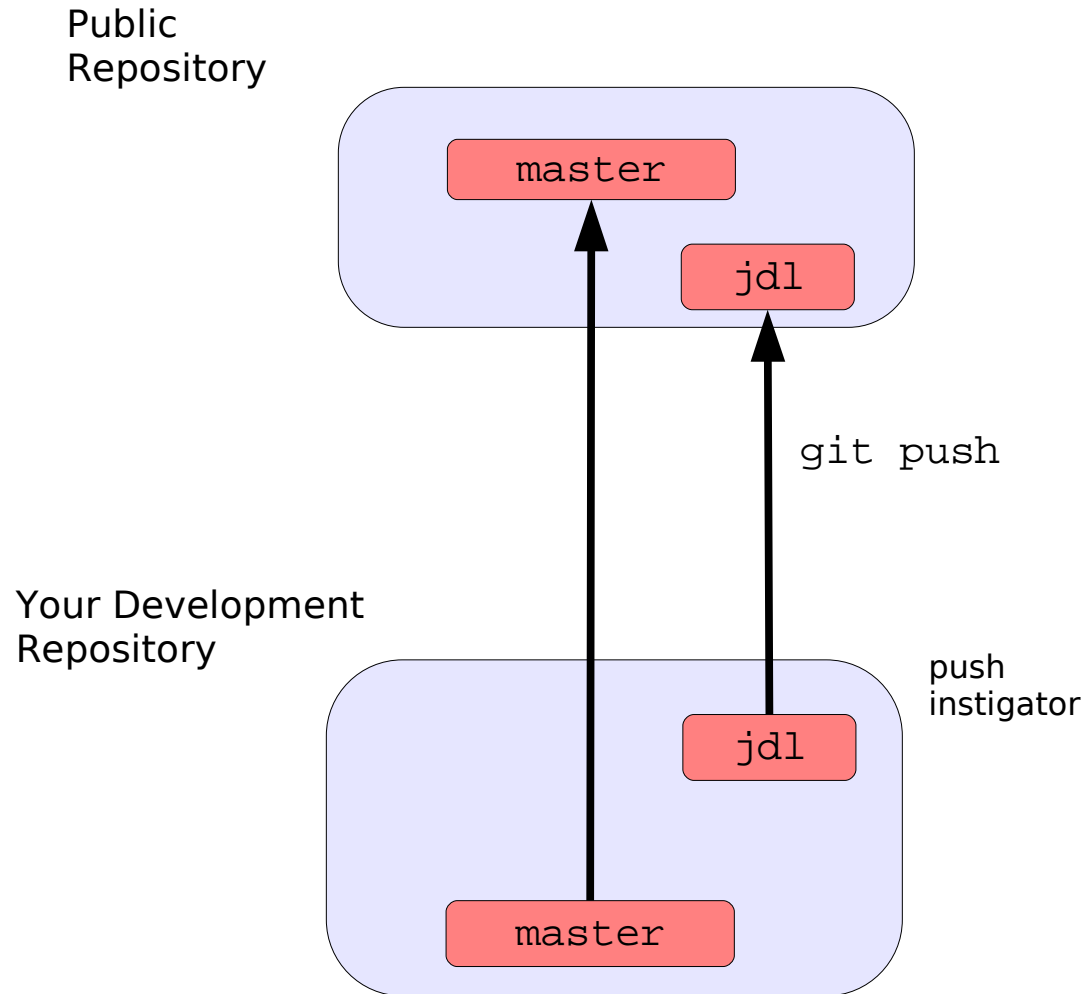
```
 100% (8/8) done
```

```
Total 8, written 8 (delta 5), reused 0 (delta 0)
```

```
Unpacking 8 objects
```

```
refs/heads/master: 38e8f8 -> a73b7d
```

Using Push: Data Flow Picture



gitk Example

The screenshot shows the gitk application window. The top menu bar includes 'File', 'Edit', 'View', and 'Help'. The main window is divided into three panes:

- Left Pane:** A commit history graph showing branches 'master' and 'pu' with various merge operations.
- Right Pane:** A list of commits with columns for author, date, and commit message. The list includes commits by Jon Loeliger, Junio C Hamano, Jakub Narebski, Peter Baumann, Timo Hirvonen, Matthias Lederhofer, Linus Torvalds, and A Large Angry SCM.
- Bottom Pane:** A detailed view of a selected commit (SHA1 ID: 5faf64cd28bf79f0c2939717d2ba117498717059). It shows the commit message, author, committer, parent, child, branch, follows, precedes, and the diff output for 'builtin-diff-files.c'.

The bottom pane also includes a 'Search' button and a 'Patch' button.